

# Native I/O backend for FS2 JVM

Name: Onur Şahin

E-mail: [sahinonur2000@hotmail.com](mailto:sahinonur2000@hotmail.com)

Address: Richard Wagner Straße 66, 79104 Freiburg, Germany

Website: <https://blog.aiono.dev/>

Github: <https://github.com/onsah/>

## Abstract

Fs2 is a “Functional, effectful, concurrent” streaming I/O library. It allows building and transforming arbitrarily complex streams possibly with side effects. With fs2, you can build streams with functional style while keeping constant memory usage and linear time complexity.

Currently, fs2 makes use of NIO in it's JVM implementation for the networking API. NIO incurs an indirection which causes non-trivial performance penalty. The project aims to implement it's I/O functionality using direct OS APIs such as `epoll/kqueue` for the JVM implementation. The Scala Native fs2 implementation already uses `epoll/kqueue`.

## How it will improve fs2?

- **Efficiency:** fs2 on JVM will perform better due to removed overhead.
- **Reduced Hardware Costs:** Because the hardware is utilized better, cheaper hardware is sufficient.
- **Tighter runtime integration:** Because I/O operations are implemented natively, runtime can be more strategical.

## Related Work

It's a part of a larger project in `cats-effect` about integrating I/O into runtime by merging compute tasks with I/O in a single thread. A detailed discussion can be found in <https://github.com/typelevel/cats-effect/discussions/3070>.

A part of the larger project, a `PollingSystem` is already implemented in `cats-effect` and it's being used in fs2 (implementation PR in `cats-effect`). Currently the `PollingSystem` for JVM is implemented using NIO Selector API. This project will implement another `PollingSystem` based on `epoll/kqueue`. So the project will be very similar to the linked PR.

## Mentors

- Arman Bilge (Website)

## Deliverables

1. Implement `epoll` based `PollingSystem` in `cats-effect`.
2. Implement networking in fs2 using `PollingSystem` API.
3. Implement `kqueue` based `PollingSystem` in `cats-effect`.
4. Performance benchmarking against the previous `PollingSystem`.
5. Documentation.

## Proposal Timeline

This project has three main parts:

- `epoll` polling system in `cats-effect`
- `kqueue` polling system in `cats-effect`
- Using `epoll/kqueue` polling systems in fs2 for networking.

I believe having one vertical slice of the implementation will help getting faster feedback and improve the overall development process. Therefore I will first implement `epoll` polling system and use it in `fs2` for networking. Then once it works properly, I will work on implementing `kqueue` polling system.

#### **May 1 - May 10:**

- Get familiar with `fs2` and `cats-effect`. Knowing how tools are used will help me see retain big picture when delving into the implementation details.
- Setup codebases locally. I already did this for `fs2` for my prior contributions.
- Keep close communication with mentors. Regularly ask questions I have regarding project design and details.

#### **May 10 - June 2**

- Investigate how networking I/O is implemented in `cats-effect` and how `NIO` is used in the runtime. Specifically study changes in this PR.
- Study how `fs2` uses `cats-effect` runtime for I/O.
- (If there is enough time) Study `epoll` and `kqueue` APIs. Possibly implement prototype programs using them.
- Study how to call system APIs from Scala. This will be necessary to utilize native I/O APIs from within `cats-effect`. Decide what technology to use. One possibility is to use `JNR`.

#### **June 2 - June 30** (Official coding period starts)

- Create a Github project under the `Typelevel` organization. Each milestone will be opened as an issue and will be linked to this project.
- Setup infrastructure to call native code from `cats-effect` in JVM implementation.
- Study `PollingSystem` API and existing JVM and Scala Native implementation.
- Implement `epoll` polling system in `cats-effect` for JVM.
- Perform automated and manual tests. Write additional automated tests if necessary.

#### **July 1 - July 16**

- Integrate `PollingSystem` for networking in `fs2` for JVM.
- Perform testing for the changes. Write additional automated tests if necessary.

#### **July 17 - August 06**

- Implement `kqueue` polling system in `cats-effect` for JVM.
- Test `fs2` with `kqueue` based `PollingSystem`.

#### **August 06 - August 15**

- Benchmark `epoll`/`kqueue` polling systems against `NIO` based polling system.
- Create a results report out of the benchmark outcomes.

#### **August 16 - August 30**

- Document new polling systems including internal implementation for their implementation details.
- Perform final testing to ensure no regressions happened and there are performance gain.
- Optimize and refactor the codebase.
- Cleanup the codebase for final submission. Ensure no commented code or useless comments exist.

### **Availability**

Generally I am available around 25-30 hours weekly during the coding period.

Currently I am in my 4th semester of my Masters. Normally I would start thesis around June, but if my proposal is accepted I will delay my thesis process so that I start working on it after GSOC. Therefore my thesis won't cause any availability issues for GSOC.

I will probably have one final exam during the coding period. To prevent any surprises I extend the standard coding period by one week in my proposal.

I have a part time job which takes around 10-15 hours weekly. At the time coding period starts, my responsibilities from it will be low because it's a university student job which is more intense during the first half of the semester. Considering I will have only one lecture and my part time job, I can comfortably give 25-30 hours weekly to this project.

## **Project Management**

### **Publishing the Code**

- I plan to open one pull request per milestone.
- Pull requests will be merged into the project specific branch first.
- Once all milestones are implemented, a final merge to the main branch will be performed.

### **Best Practices**

- I will follow project's contribution rules such as formatting, and certain coding conventions.
- I will clarify anything uncertain by discussing them with my mentor and community.

### **Sharing**

- During the project, I plan to write about it in my blog so that I can share the knowledge I gained with others. Also it will help me to mentally organize the project and serve as a documentation.

## **My Background**

I have a keen interest in systems programming and functional programming which I demonstrate with my toy programming language implementation. I have a work background where I used Scala Futures professionally so concepts like blocking/nonblocking I/O, event loops and thread pools are familiar to me. My professional Scala experience is around 2 years so I am fairly comfortable with the language. Though I don't have much experience on Cats before I got into the community, I am working on it by following starter guides and I will continue to familiarize myself.

Before submitting this proposal, I introduced myself to Typelevel project to make early contributions and show my interest in the project. Thankfully, Arman Bilge had was very attentive and helpful. With his guidance, I worked on implementing non-blocking process API in fs2. I have successfully landed changes into fs2 with <https://github.com/typelevel/fs2/pull/3539> and <https://github.com/typelevel/fs2/pull/3548>.