

Native I/O backend for FS2 JVM

Name: Onur Şahin

E-mail: sahinonur2000@hotmail.com

Address: Richard Wagner Straße 66, 79104 Freiburg, Germany

Website: <https://blog.aiono.dev/>

Github: <https://github.com/onsah/>

Abstract

`fs2` is a “Functional, effectful, concurrent” streaming I/O library. It allows building and transforming arbitrarily complex streams possibly with side effects. With `fs2`, you can build streams with functional style while keeping constant memory usage and linear time complexity.

Currently, `fs2` makes use of NIO in its JVM implementation for the networking API. NIO uses a lot of synchronization (locking) and causes indirection which result with non-trivial performance penalty. The project aims to implement its I/O functionality using direct OS APIs such as `epoll/kqueue` for the JVM implementation. The goal is to reduce locking and indirection to increase the performance. The Scala Native `fs2` implementation already uses `epoll/kqueue`.

How it will improve fs2?

- **Efficiency:** `fs2` on JVM will perform better due to removed overhead.
- **Reduced Hardware Costs:** Because the hardware is utilized better, cheaper hardware is sufficient.
- **Tighter runtime integration:** Because I/O operations are implemented natively, runtime can be more strategic.

Related Work

It's a part of a larger project in `cats-effect` about integrating I/O into runtime by merging compute tasks with I/O in a single thread. A detailed discussion can be found in <https://github.com/typelevel/cats-effect/discussions/3070>.

A part of the larger project, a `PollingSystem` is already implemented in `cats-effect` and it's being used in `fs2` ([implementation PR in cats-effect](#)). Currently the `PollingSystem` for JVM is implemented using NIO Selector API. This project will implement another `PollingSystem` based on `epoll/kqueue`. So the project will be very similar to the linked PR.

Mentors

- Arman Bilge ([Website](#))
- Antonio Jimenez ([Github](#))

Deliverables

1. `epoll` based `PollingSystem` in `cats-effect`.
3. `kqueue` based `PollingSystem` in `cats-effect`.
2. `fs2` migration to `PollingSystem`.
4. Performance benchmarking against the previous `PollingSystem`.
5. Documentation: Both internal and external.

Proposal Timeline

This project has three main parts:

- `epoll` polling system in `cats-effect`
- `kqueue` polling system in `cats-effect`
- Using `epoll/kqueue` polling systems in `fs2` for networking.

I believe having one vertical slice of the implementation will help getting faster feedback and improve the overall development process. Therefore I will first implement `epoll` polling system and use it in some parts of `fs2`. Then once it works properly, I will work on implementing `kqueue` polling system. Then I will work on remaining places to migrate in `fs2`.

Mentors informed me that, migrating all appropriate modules in `fs2` into `PollingSystem` might not be realistic for the project duration. Therefore, I added some modules into the timeline as optional.

May 1 - May 10:

- Get familiar with `fs2` and `cats-effect`. Knowing how tools are used will help me see retain big picture when delving into the implementation details.
- Setup codebases locally. I already did this for `fs2` when I previously contributed.
- Keep close communication with mentors. Regularly ask questions I have regarding project design and details.

May 10 - June 2

- Investigate how networking I/O is implemented in `cats-effect` and how `NIO` is used in the runtime ([this PR](#)).
- Study existing migration to `PollingSystem` in `fs2` ([this PR](#)).
- Study `epoll` and `kqueue` APIs. Possibly implement prototype programs using them.
- Study how to call system APIs from Scala. This will be necessary to utilize native I/O APIs from within `cats-effect`. Decide what technology to use. One possibility is to use [JNR](#).

June 2 - June 6 (Official coding period starts)

- Create a Github project under the Typelevel organization. Each milestone will be opened as an issue and will be linked to this project.
- Setup infrastructure to call native code from `cats-effect` in JVM implementation.

June 7 - June 29

- Implement `epoll` polling system in `cats-effect` for JVM.
- Perform automated and manual tests. Write additional automated tests if necessary.

June 30 - July 27

- Migrate `SocketGroup` / `Socket` in `fs2` to use `PollingSystem`.
- Migrate `UnixSockets` in `fs2` to use `PollingSystem`.
- Ensure proper documentation is in place.
- Test that `SocketGroup`, `Socket` and `UnixSockets` modules work correctly with `epoll` `PollingSystem`. Write automated tests.
- (Optional) if there is enough time, also migrate `DatagramSockets` and `Process` modules.

July 28 - August 10

- Implement `kqueue` polling system in `cats-effect` for JVM.
- Test that `SocketGroup`, `Socket` and `UnixSockets` modules work correctly with `kqueue` `PollingSystem`. Write automated tests.

August 11 - August 17

- Benchmark `epoll`/`kqueue` polling systems against `NIO` based polling system.
- Create a results report out of the benchmark outcomes.

August 18 - August 30

- Document new polling systems including internal implementation for their implementation details.

- Perform final testing to ensure no regressions happened and there are performance gain.
- Optimize and refactor the codebase.
- Cleanup the codebase for final submission. Ensure no commented code or useless comments exist.

Availability

Generally I am available around 25-30 hours weekly during the coding period.

Currently I am in my 4th semester of my Masters. Normally I would start thesis around June, but if my proposal is accepted I will delay my thesis process so that I start working on it after GSOC. Therefore my thesis won't cause any availability issues for GSOC.

I will probably have one final exam during the coding period. To prevent any surprises I extend the standard coding period by one week in my proposal.

I have a part time job which takes around 10-15 hours weekly. At the time coding period starts, my responsibilities from it will be low because it's a university student job which is more intense during the first half of the semester. Considering I will have only one lecture and my part time job, I can comfortably give 25-30 hours weekly to this project.

Project Management

Publishing the Code

- I plan to open one pull request per milestone.
- Pull requests will be merged into the project specific branch first.
- Once all milestones are implemented, a final merge to the main branch will be performed.

Best Practices

- I will follow project's contribution rules such as formatting, and certain coding conventions.
- I will clarify anything uncertain by discussing them with my mentor and community.

Sharing

- During the project, I plan to write about it in my [blog](#) so that I can share the knowledge I gained with others. Also it will help me to mentally organize the project and serve as a documentation.

My Background

I am passionate about systems programming and functional programming, with a strong interest in building scalable and efficient systems. As a demonstration of this, I developed `Flux_rs`, a toy programming language implementation (with bytecode VM) that deepened my understanding of language design and runtime mechanics.

Open source has a very special for me since I started programming, and I actively contribute to projects that align with my interests in systems programming and developer tools. Some highlights include:

- Maintained a [Gnome Shell Extension](#) with 16 stars
- [Contributed to Nu Shell's standard library](#), enhancing usability for a growing community.
- Making [multiple contributions to Jakt typechecker](#).

Professionally, I have 2 years of experience using Scala, including working with Scala Futures, thread pools, and non-blocking I/O in production environments. This hands-on experience has given me a solid grasp of concurrency models and performance considerations in real-world systems.

While I'm still deepening my familiarity with Cats, I've been actively learning the ecosystem through starter guides and practical contributions. To demonstrate my commitment to the Typelevel

ecosystem, I engaged early by contributing to fs2. Under the mentorship of Arman Bilge, I successfully implemented a non-blocking process API, with my changes merged in <https://github.com/typelevel/fs2/pull/3539> and <https://github.com/typelevel/fs2/pull/3548>. This experience not only strengthened my understanding of functional effect systems but also showcased my ability to deliver meaningful contributions to open-source projects.

I'm excited to continue contributing to Typelevel and further grow my expertise in functional programming through GSoC.